Chapter 1

IBCM: The Itty Bitty Computing Machine

The child receives data through the sense organs; the child also has some inborn processing capacities –
otherwise it would not be able to learn – but in addition, some "information" or "programs" are built-in
at birth ... there is a working memory ... and there is a permanent memory ... so there must be some
inner "language" or medium of representation ... Jerry Fodor ... has discussed this inbuilt "language
of thought," which is similar conceptually to the "machine language" that is built into the personal
computer.
    – John Cleverly, in "Visions of childhood: Influential models from Locke to Spock" [3]

1.1   Introduction

Machine language, or **machine code**, is the set of instructions that a computer's central processing unit
(CPU) understands. It is the binary 0's and 1's that form instructions for the CPU to execute. When we
compile a program, that program is eventually converted to binary machine code, which we can then run.
Each different CPU has a different machine language that it understands, although CPU families tend to
understand a very similar language.

Programming in machine language requires one to write the code in hexadecimal notation, manually
encoding the instructions one at a time. For this reason, machine language can often be difficult to program
in. This is especially true with the complexities of modern instruction sets on processor families
such as the x86 and MIPS.

One will often write a program in **assembly language**, which is a (somewhat) higher-level manner to
create a program. In assembly language, one can specify to add two values together through a command
such as `sub esp, 24`, and need not write it in the hexadecimal notation of 0x83ec20. An **assembler** is
a program that converts assembly language into machine language. In fact, compilers typically produce
assembly language, and then call the assembler to produce the machine code.

Learning how machine language works allows us to truly understand how a machine operates at a
very low level.

In this chapter we introduce a simplified machine language called the Itty Bitty Computing Machine,
or IBCM. This machine language is designed to be simple enough that it can be learned in a reasonable
amount of time, but complex enough that one can use it to write a wide range of programs. It is intended
to teach many of the important concepts of a modern machine language without having to deal with the

1

complexities of – and time required to learn – modern CPU machine languages.

1.2   Memory Hierarchy

Typically, computer programs do not differentiate between the various levels of memory. Programs tend to view memory as a monolithic whole, and allocate memory as needed. In fact, there are many different levels of memory. Fast memory – formally called *static random-access memory* or SRAM – is expensive to make, and computers would be quite expensive if all memory were fast memory. In fact, some super computers did just this – the Cray Y-MP, a super computer from the late 1980's, used only SRAM, and it cost approximately 10 million dollars.

Instead, most computers use primarily dynamic random-access memory, or DRAM, as the primary component of main memory, and use a smaller amount of SRAM to speed up computations. SRAM is used in the *cache*, although we will not go into detail about caches here.

The four primary levels of memory are listed below. Cost decreases and storage capacity increases as one moves down the list.

- CPU registers
- Static random-access memory (SRAM) in various levels of cache (L1 cache, L2 cache, etc.)
- Dynamic random-access memory (DRAM) used in main memory
- Hard drive storage, either on a traditional hard disk with rotating platters, or on solid state hardware.

This chapter will largely ignore storing data in cache or on the hard drive, as that is typically managed by the operating system. Instead, we will primarily focus on the values kept in the registers in a CPU, and in main memory.

Figure 1.1: Memory Hierarchy

Main memory is relatively slow, compared with the speed of a CPU. In addition, many CPUs can only allow one memory access per instruction. Thus, the other values must be kept in a location that is not main memory: the CPU's registers. A register is a small area of memory in the CPU where intermediate computation values can be stored. A modern CPU will have registers that each are 32 or 64 bits in size, and may have 12-32 of these registers.

Thus, there are instructions – in both machine language and assembly language – that move data between the CPU's registers and main memory. We will see these instructions shortly.

1.3   IBCM Principles of Operation

1.3.1   Machine Description

The Itty Bitty Computing Machine (IBCM) is a very simple computer. In fact, it's so simple that no one in their right mind would build it using today's technology. Nevertheless – except for limits on problem size – any computation that can be performed on the most modern, sophisticated computer can also be performed on the IBCM. Its main virtues are that it can be taught quickly and provides context for talking about more recent architectures.

The IBCM contains a single register, called the accumulator. Most early computers had just an accumulator, and many current micro-controllers are still accumulator machines. IBCM's accumulator can hold a 16-bit 2's complement integer. In addition, the IBCM has two other registers: the instruction register (IR), which contains the current instruction being executed, and the program counter (PC), which contains the address of the next instruction to be executed. However, we will not directly use those two registers.

The IBCM has 4,096 addressable memory locations and each of these locations holds 2 bytes (16 bits). Thus, the IBCM technically has 8,192 bytes (8 Kb) of memory.

The IBCM handles input by reading in single values (either a hexadecimal value or an ASCII character) into the accumulator. Similarly, output is handled by writing the value in the accumulator as either a hexadecimal value or an ASCII character.

1.3.2   Instructions

The IBCM equivalent of "statements" in a higher level language are very simple instructions. Each instruction is encoded in a 16-bit word in one of the formats shown in Figure 1.2. Shaded areas in the figure represent portions of the instruction word whose value does not matter. For example, a word whose leftmost four bits are zero is an instruction to halt the computer.

bit:	15	14	13	12	11 10	...	0	
	0	0	0	0	(unused)			halt
	0	0	0	1	I/O op	(unused)		I/O
	0	0	1	0	shift op	(unused)	count	shifts
	opcode				address			others

Figure 1.2: IBCM instruction format

A note on using hexadecimal notation: IBCM is a binary computer. Internally all operations are performed on 16-bit binary values. However, because it's so tedious and error-prone for humans to write or read 16-bit quantities, all of IBCM's input/output is done either as ASCII characters or 4-digit hexadecimal numbers. Remember that these are just external shorthands for the internal 16-bit values!

Also, hexadecimal values are often prefixed by a '0x', such as '0x1234'.

Halt

Any instruction where the opcode is zero (i.e., the first 4 bits are all zero) will halt the IBCM. It does not matter what the remaining 12 bits are. Not much else to say on this one.

Input and Output

The input and output (or 'I/O') instructions move data between the accumulator and the computer 'devices' – the keyboard and screen. Data can be moved either as hexadecimal numbers or as an ASCII character; in the later case, only the bottom (least significant) 8 bits of the accumulator are involved. The next two bits after the opcode specify if the instruction is an input or output, and if it will use hexadecimal values or ASCII characters. The four possibilities (in/out, hex/ASCII) that are specified by the bits 11 and 10 of the instruction word as shown in Table 1.1. Being as there are only four possibilities, the full hexadecimal encoding of each of the four I/O instructions is also listed.

bit 11	bit 10	operation	hex value
0	0	read a hexadecimal word (four digits) into the accumulator	0x1000
0	1	read an ASCII character into the accumulator bits 8-15	0x1400
1	0	write a hexadecimal word (four digits) from the accumulator	0x1800
1	1	write an ASCII character from the accumulator bits 8-15	0x1c00

Table 1.1: IBCM input/output bit values

Shifts and Rotates

Shifting and rotating are common operations in computers. Shifting means moving data to the right or left; rotating is much the same thing except that bits that "fall off" one end are reinserted at the other end.

The examples below assume that the value being rotated is one byte (8 bits) in length. In reality, the IBCM performs shifts and rotates on the 16 bit value in the accumulator. We choose to describe these operations with 8 bit quantities to make the explanation easier to understand.

Diagrammatically, a "right shift" of "n positions" looks like Figure 1.3 when $n = 3$:


Figure 1.3: A right shift of 3

Each bit is moved n positions to the right (three in this case). Alternatively, it is moved one bit to the right n times. This causes the original n right-most bits to fall off the right end and n new (zero) bits to be inserted at the left end. A "left shift" is much the same, except that bits move to the left, as shown in Figure 1.4.

Original

0	1	0	0	1	1	1	0

Shifted 0 1 0

| 0	1	1	1	0	0	0	0

Figure 1.4: A left shift of 3

As noted previously, rotates are like shifts except that the bits that fall off one end but are reinserted at the other end. Figure 1.5 is a left rotation; right rotates are left to your imagination.

Original

0	1	0	0	1	1	1	0

Shifted

| 0	1	1	1	0	0	1	0

Figure 1.5: A left rotation of 3

Note that the shift instructions uses bits 11-10 to specify the shift operation. The first bit specifies if it is a shift (0) or rotate (1), and the second bit specifies if it is moving left (0) or right (1). This is shown in Table 1.2.

bit 11	bit 10	operation
0	0	shift left
0	1	shift right
1	0	rotate left
1	1	rotate right

Table 1.2: IBCM shift/rotate bit values

In addition, bits 3-0 of the shift instructions specify the "shift count" – that is, the number of bits positions that the data is to be shifted (or rotated).

A left rotate of 3, which is what is shown in Figure 1.5, would have the opcode set to 2 (binary 0010), the shift/rotate bit set to 1, the left/right bit set to 0, and the shift count set to 3. The encoded instruction, in binary, is shown in Figure 1.6. Note that the grayed-out part of the table are the bits whose value does not matter; we set them to 0.

0 0 1 0	1	0	0 0 0 0 0 0	0 0 1 1	= 0x2803
opcode	rot. bit	left bit	unused bits	shift count	

Figure 1.6: IBCM instruction for a right rotate of 3

As shown in the table, the hexadecimal encoding of the instruction is 0x2803.

Other Instructions

All other bit combinations in the left-most four bits either specify an arithmetic instruction or a control instruction – similar to assignment statements and `gotos` in a high level language. The first four bits are called the "op" field of the instruction; the value of this field is often called the "opcode". The "address" portion of the instruction generally specifies an address in memory where an operand (variable) will be found.

There are 13 IBCM operations of this form – eight that manipulate data and five that perform control. The data manipulation operations all involve the "accumulator" and typically data from a memory location is specified by the address portion of the instruction (the `not` and `nop` instructions are the only two different ones). The result of data manipulation operations is recorded in the accumulator. Thus, the "add" instruction forms the arithmetic sum of the present contents of the accumulator with the contents of the memory location specified by "address" and puts the result back into the accumulator. Thus, it is similar to the primitive assignment statement "`accumulator = accumulator + memory[address]`".

The control instructions determine the next instruction to be executed. The "jump" instruction, for example, causes the next instruction executed to be the one at the location contained in its address field. If you think of the address of a memory cell like a label in a high level language, then jump is just "`goto address`".

Two of the control instructions are conditional; they either cause a change in the control flow or not, depending on the value of the accumulator. The simplest of the control instruction is `nop`; it does nothing.

Table 1.3 describes the function of each of these 13 IBCM instructions; both English and programming language-like explanations are given for each instruction. In the latter, "a" is the accumulator, "addr" is the value of the address portion of the instruction, and "mem[]" is memory.

op	name	HLL-like meaning	English explanation
3_{16}	load	a := mem[addr]	load accumulator from memory
4_{16}	store	mem[addr] := a	store accumulator into memory
5_{16}	add	a := a + mem[addr]	add memory to accumulator
6_{16}	sub	a := a – mem[addr]	subtract memory from accumulator
7_{16}	and	a := a & mem[addr]	logical 'and' memory into accumulator
8_{16}	or	a := a | mem[addr]	logical 'or' memory into accumulator
9_{16}	xor	a := a $\oplus$ mem[addr]	logical 'xor' memory into accumulator
A_{16}	not	a := $\sim$ a	logical complement of accumulator
B_{16}	nop		do nothing (no operation)
C_{16}	jmp	goto 'addr'	jump to 'addr'
D_{16}	jmpe	if a = 0 goto addr	jump to 'addr' if accumulator equals zero
E_{16}	jmpl	if a < 0 goto addr	jump to 'addr' if accumulator less than zero
F_{16}	brl	branch and link	jump (branch) to 'addr'; set accumulator to the value of the the "return address" (i.e., the instruction just *after* the brl)

Table 1.3: IBCM opcodes

A few words of additional explanation are necessary for some of these operations, all of which are fairly standard for most computers.

1. Arithmetic operations may "overflow" or "underflow"; that is, the magnitude of the result may be larger than can be represented in 16 bits. The programmer is responsible for ensuring that this

doesn't happen.

2. The logical operations (`and`, `or`, `xor`, and `not`) perform bit-wise operations on the operands. The Boolean operations themselves are shown in Figure 1.7.

and	0	1
0	0	0
1	0	1

or	0	1
0	0	1
1	1	1

xor	0	1
0	0	1
1	1	0

not	
0	1
1	0

Figure 1.7: Boolean operation reference

3. The `not` instruction only inverts the bits in the accumulator, and does not use a memory location; the 'address' part of the instruction is ignored.

4. Likewise, the `nop` instruction ignores the 'address' part of the instruction.

5. The branch and link instruction is used for subroutine calls, as discussed later.

1.3.3 Other opcodes

When writing an IBCM program, one will typically write out the IBCM opcodes, such as `add` and `store`. There are a few other opcodes that are not listed above, but that will often appear in an IBCM file:

- `dw` (for "declare word"), for declaring variables
- `readH` and `printH` are for reading or writing a hexadecimal value
- `readC` and `printC` are for reading or writing an ASCII character
- `shiftL` and `shiftR` are for the shifts
- `rotL` and `rotR` are for the rotations

1.4   Sample Program

Consider the IBCM program shown in Listing 1.1. This program does not compute a useful result; that's coming next. Instead, it is intended to show how the IBCM works.

Address	Instruction	Opcode	Address
000	3000	load	000
001	5000	add	000
002	6001	sub	001
003	8003	or	003
004	a000	not	N/A
005	4000	store	000
006	f000	brl	000

Listing 1.1: Sample IBCM program

Let's trace what this program does. All the values below are in hexadecimal. All addresses are represented using three digits.

1. Address 000: Instruction value 3000. Opcode 3 is a `load`, with an address of 000. This will load the value in memory at address 000 into the accumulator. The value in address 000 is 3000 – it is both the instruction being executed and the data being loaded. The accumulator is now 3000.

2. Address 001: Instruction value 5000. Opcode 5 is an `add`, with an address of 000. This will add the value in memory at address 000 to the accumulator, and store the result in the accumulator. The value in address 000 is 3000, so the result (and the new accumulator value) is 6000.

3. Address 002: Instruction value 6001. Opcode 6 is a `sub`, with an address of 001. This will subtract the value in memory at address 001 from the accumulator, and store the result in the accumulator. The value at address 001 is 5000; 6000-5000=1000. Thus, the new accumulator value is 1000.

4. Address 003: Instruction value 8003. Opcode 8 is an `or`, with an address of 003. This will perform a bit-wise or'ing of the value in memory address 003 with the value in the accumulator, and store the result in the accumulator. The value in address 003 is 8003 – again, we are using the same value for the instruction being executed and for the data being used.

   To perform a bit-wise logical operation, write out the full bit values for each of the operands, and perform the bit-wise operation (or, in this case) on each of the bits in each column.

   $$
   \begin{array}{rclcccc}
   0x8003 & = & 1000 & 0000 & 0000 & 0011 \\
   \lor \quad 0x1000 & = & 0001 & 0000 & 0000 & 0000 \\
   \hline
   & & 1001 & 0000 & 0000 & 0011 & = \quad 0x9003
   \end{array}
   $$

   Thus, the value in the accumulator is 9003.

5. Address 004: Instruction value a000. Opcode a is a `not` operation. This opcode ignores the address portion of the instruction, as the operation just inverts all the bits of the accumulator. The bit value of the accumulator prior to the not operation is listed in the previous step.

   $$
   \begin{array}{rclcccc}
   \lnot \quad 0x9003 & = & 1001 & 0000 & 0000 & 0011 \\
   \hline
   & & 0110 & 1111 & 1111 & 1100 & = \quad 0x6ffc
   \end{array}
   $$

   Thus, the value in the accumulator is 6ffc.

6. Address 005: Instruction value 4000. Opcode 4 is a `store` instruction, with an address of 000. This will store the current value in the accumulator (6ffc) into memory at address 000. The previous value in that spot (3000) is, of course, overwritten.

7. Address 006: Instruction value f000. Opcode f is a `brl` (branch and link) instruction, with address 000. This will store the address of the next instruction (i.e., the address after 006, or 007) in the accumulator, and jump to the specified address of 000. Thus, the value in the accumulator after this instruction is 0007, and the next instruction to be executed is address 000.

8. Address 000: Instruction value 6ffc. Note that this value was written to memory two steps prior (the instruction at address 005), and control jumped to this instruction from the previous instruction (the instruction at address 006). Opcode 6 is a `sub`, with address ffc. As all values in IBCM's memory are initialized to zero, the value in address ffc is thus zero. Thus will subtract zero from the accumulator yielding the original value of 0007, which is what is stored in the accumulator after this instruction is executed.

   Program control will continue. The next time the accumulator executes the instruction at address 005 (instruction value 4000; opcode 4 is a `store`), the value written to address 000 will be 5ffc (opcode 5 is an

add). The next time around, it will write 6ffc (opcode 6 is a sub). Thus, this program will loop forever, alternately writing 6ffc and 5ffc to address 000 at the end of each loop.

One of the important points to note in the above program is that the distinction between data and instructions is blurred. The value 6ffc is computed through a series of instructions, and then stored – as data – at address 000. But when the program control reaches that point, the value of 6ffc – that was previously data – now becomes an instruction.

We will see additional examples of using data as instructions, and using instructions as data, in the example programs section, below.

1.5 The IBCM Simulator/Debugger

To run an IBCM program, you can view the online simulator at http://libra.cs. virginia.edu/ibcm [2]. The instructions listed in this section are also listed at that website. An image of the online emulator is shown in Figure 1.8.

The simulator reads in text files, and proceeds to simulate the result of an IBCM program execution.

To load a file, use the Browse button at the top of the simulator page. Find the IBCM file, and click on the "Load" button. The format of your program file is very rigid – the first four characters of each line are interpreted as a hexadecimal number. The number on the first line is loaded into location zero, the next into location one and so on. Characters after the first four on each line are ignored, so you should use them to comment the code. Blank lines are not allowed, nor are lines that do not start with four hexadecimal digits (i.e., no 'comment' lines). An invalid file will either not load up at all, or will load up gibberish.

The left side of the simulator lists all the memory locations (using the hexadecimal address), the value in memory (if any), and the PC (program counter) value. Note that any blank value field is interpreted as'0000', as per the IBCM specification. When first loaded, the simulator may uninitialized values blank to increase readability.

Figure 1.8: Web-based IBCM interface

The value column consists of a series of text boxes, which allow you to directly edit the values in memory. The simulator will read the current memory location from the appropriate text box when executing an instruction. You can undo any edits by using the Revert button, described below. The IBCM simulator does not check to ensure that your entered values are valid – it is up to the user to do this. Note that all hexadecimal values must be 4 digits (i.e. '0000', not '0'), or else the simulator may not function correctly. Be sure to read the section about crashing browsers and losing your work, below. There is no way to save

your edited work – you will need to copy the changes by hand. This is party due to a browser limitation, and partly due to the fact that memory editing is meant to be a debugging tool, not a means to write entire IBCM programs from scratch.

The 'PC' column lists the current value of the program counter. It can have three values. Normally, it will have a left pointing arrow ('<-'), which indicates the next instruction that will be executed. If the simulator is waiting for input, it will have a capital 'I' (for Input) next to the input instruction that is currently awaiting a value – and the "Input" text on the right side will blink. Lastly, if the program has halted, then a capital 'H' (for Halt) will be displayed next to the halt instruction that was executed.

On the right side, the values of the accumulator and program counter are listed, both in hexadecimal notation. As mentioned above, the PC field will also display, on the left side next to the hexadecimal address, if the simulator is awaiting input or is halted. The Input box is used to read in user input when a program requests it. When the simulator is waiting for input, it will flash the 'Input' text. In addition, the simulator will specify which type of input is being requested: 'hex' or 'asc' for hexadecimal or ASCII input, respectively. Note that for entering a hexadecimal value, you do not need to enter all 4 digits: i.e., you can enter '12' instead of '0012'. This is distinctly different that editing memory locations (you have to enter all 4 digits for those). If you enter multiple characters for ASCII input, it will only read in the first one.

Below this are four buttons. Two control execution: Run, which will start a program executing, and Step, which will execute a single instruction. Note that Run will execute until either a halt command is reached, or until an input command is reached. The other two buttons control the resetting of the IBCM program. The Reset button will reset the PC and accumulator, in effect allowing the program to run again. It will not, however, modify any memory locations. The Revert button will do what a Reset does, but will also revert the memory locations to what they were when the file was last loaded; it does not load the file from disk again. Thus, if you have edited any of the memory locations (or the IBCM program has modified them), then those changes will be erased on a Revert, but not on a Reset. Note that a Revert will not modify memory addresses outside the range that was loaded (i.e. any 'blank' values).

Any output is displayed in the text area below these buttons. Each output command prints the value (hexadecimal or ASCII character) on a separate line.

A few notes:

- You will notice a slight delay when loading the simulator page. This is due to the fact that a number of scripts are run when the page loads (to initialize the memory table of 4096 elements, for example), and this takes a bit of time. How long this takes is determined by how fast a computer it is running on, as the scripts are run on the client side.

- Upon entering input, hitting Enter is considered the same as hitting the Run button again. If you only want to execute a single instruction after an input, you must click on the Step button.

- The simulator does minimal error checking with the input from the keyboard during program execution – it is the user's responsibility to ensure that the input is properly formatted. The only error checking that is done is to ensure that a non-empty string was entered (if it was, then the simulator waits for more input).

- Because of the limitations of threads running in web browsers, there is no way to terminate a program that is stuck in an infinite (or very long) loop – the browser will not allow a polling (checking) to see if a Stop button was pressed, for example. Some browsers will pause the script after a minute

or so of execution, and ask if the user wants to continue. Alternatively, you can close the web browser and restart. Note that this means if you have edited any of the memory locations, and your browser hangs or is restarted, you will lose any and all changes you have made to the memory locations!

- The simulator page is a PHP script, which means that it will not work if you are viewing it as a local file (if the beginning of your URL is "file://" instead of "http://"), or if the web server hosting this page does not have PHP installed (this latter restriction includes UVa's Collab web server).

- Browser compatibility: It has been tested in Internet Explorer under Windows, Safari under Mac OS X, and Firefox under Windows, Mac OS X, and Ubuntu Linux. Note that the display of the changing of the values as the simulation is run (the PC, memory values, etc.) will only work on some browser / operating system combinations (Firefox, in particular, works well for this). The other browsers will have the same end state after the program is run, but will not animate the execution of the IBCM program when 'Run' is pressed (the 'Step' command will still animate each step).

- The format of your program file is very rigid – the first four characters of each line are interpreted as a hexadecimal number. The number on the first line is loaded into location zero, the next into location one and so on. Characters after the first four on each line are ignored, so you should use them to comment the code; example code will be discussed in class.

- When you execute a `halt`, the simulator will halt with the PC pointing at the `halt`.

1.6 Writing IBCM Programs

1.6.1 Complex control structures

The control structures that we are familiar with in higher level languages can be implemented in IBCM, albeit with a bit more effort. Consider a typical pseudo code if-then-else conditional shown in Listing 1.2 to the right.

```
if ( B == 0 )
  S1;
else
S2
```

Listing 1.2: **if** pseudo code

The IBCM code for this conditional is shown in Listing 1.3. Because we do not know many of the required addresses (of variable B, or where S1 and S2 start in memory), we have chosen to leave it in assembly format (i.e., using opcodes) rather than provide machine code.

If we wanted to compare B to a different value, such as 5, we would have to load B into the accumulator, subtract 5 from that, and then perform a `jmpe` or `jmpl`. This is illustrated further below, when discussing the while loop.

```
load B
jmpe S1
S2: ...
jmp done;
S1: ...
done: ...
```

Listing 1.3: IBCM **if** code

Presumably, the ellipses at labels S1 and S2 would have some set of IBCM opcodes to execute.

Loops are also easily converted to IBCM. Note that a `for` loop is just a `while` loop, but with a statement performed before the loop starts (the "for init"), and a statement performed at the end of each loop iteration (the "for update"). Consider a straight forward while loop shown in Listing 1.4.

```
while ( B >= 5 )
  S;
```

Listing 1.4: **while** pseudo code

We cannot directly compare a variable to the value 5; we can only compare it to zero via the `jmpe` and `jmpl` instructions. Our IBCM code is shown in Listing 1.5.

If $B < 5$, then we want our loop to terminate. If $B < 5$, then $B - 5 < 0$, so we will execute a `jmpl` to break out of the loop. If $B = 5$, then $B - 5 = 0$, and we will continue in the loop.

```
loop:  load B
sub  five
jmpl  done
S:  ...
jmp  loop
done:  ...
```

Listing 1.5: IBCM **while** code

1.6.2   General tips

When writing IBCM programs, we recommend these steps:

1. Write the pseudo code first – or even actual code – to make sure your algorithm works as desired. If you have a bug in the design of your algorithm, then you will never get your IBCM code to work.

2. Write the program using IBCM opcodes, so that it looks like assembly: `add one`, `store x`, etc. Comment this clearly!

3. Trace this assembly code, by hand, from beginning to end. It will be *far* easier to find a bug in your program in the assembly stage than in the hexadecimal code stage.

4. Finally, translate it to machine code, and run it in the simulator.

We cannot stress enough how important it is to first make sure that the algorithm works, then to write and trace the assembly versions of the program. Debugging hexadecimal machine code is not much fun.

1.7   Example Programs

We present a number of IBCM example programs to help you get acquainted with the language.

1.7.1   Summation

The first program will compute the sum of the integers 1 through n, where n is read from the keyboard; the resulting sum is printed to the screen. The program then halts after printing the sum. The C++ code for the program is shown in Listing 1.6 – this is presented to help show the conversion to IBCM.

```cpp
int main(void) {
  int n, s = 0;
  cin >> n;
  for ( int i = 0; i <= n; i++ )
    s += i;
  cout << s << endl;
}
```

Listing 1.6: C++ summation program

The IBCM program is shown in Listing 1.7. Note that the only part of the program that the simulator reads in is the first four characters on the line. The rest of the line in the text file is solely for comments. The column headers shown in the figure are heavily abbreviated to fit in the width of a column, but are, in order: the actual 4 hexadecimal digit memory value, the hexadecimal location (used for determining jump targets and variable addresses), the label (used to refer to jump and variable targets), the opcode (from

Table 1.3), the target address (which refers to a given label), and any English comments. Note that the column headers would not be in the input file; only the IBCM hexadecimal instructions. They are included in the listing for clarity.

Mem	Loc'n	Label	Opcode}	Addr	Comments
C006	00		jmp	init	jmp past vars
0000	01	i	dw		int i
0000	02	s	dw		int s
0000	03	n	dw		int n
0001	04	one	dw		
0000	05	zero	dw		
1000	06	init	readH		read n
4003	07		store	n	
3004	08		load	one	i = 1
4001	09		store	i	
3005	0A		load	zero	s = 0
4002	0B		store	s	
3003	0C	loop	load	n	if i>n, jmp xit
6001	0D		sub	i	
E016	0E		jmpl	xit	
3002	0F		load	s	s += i
5001	10		add	i	
4002	11		store	s	
3001	12		load	i	i += 1
5004	13		add	one	
4001	14		store	i	
C00C	15		jmp	loop	goto loop
3002	16	xit	load	s	print s
1800	17		printH		
0000	18		halt		halt

Listing 1.7: IBCM summation program

1.7.2 Array usage

For this program, we will compute the sum of the elements of an array, print this sum on the screen, and then halt. Note that the array is just a series of sequential spots in memory; it could be part of the program itself, or a series of data values after the program itself. The address of the first element of the array and the size of the array are to be read from the keyboard.

The pseudo code for this program is shown in Listing 1.8. C++ does not easily allow for reading in an array address (while technically possible, it is not very practical), so a C++ version of this program is not as

```
read A
read N
s = 0
i = 0
while (i < N)
        s += a[i]
        i += 1
print s;
```

Listing 1.8: Array index pseudo code

useful as a pseudo code version.

The IBCM code is shown in Listing 1.9.

Mem	Loc'n	Label	Opcode	Addr	Comments}
C00A	00		jmp	start	skip around the variables
0000	01	i	dw	0	int i
0000	02	s	dw	0	int s
0000	03	a	dw	0	int a[]
0000	04	n	dw	0	
0000	05	zero	dw	0	
0001	06	one	dw	1	
5000	07	adit	dw	5000	
0000	08				leave space for changes
0000	09				
1000	0A	start	readH		read array address
4003	0B		store	a	
1000	0C		readH		read array size
4004	0D		store	n	i = 0; s = 0
3005	0E		load	zero	
4001	0F		store	i	
4002	10		store	s	
3004	11	loop	load	n	if (i >= N) goto xit
6001	12		sub	i	
E020	13		jmpl	xit	
D020	14		jmpe	xit	
3007	15		load	adit	form the instruction to add a[i]
5003	16		add	a	
5001	17		add	i	
401A	18		store	doit	plant the inst into the program
3002	19		load	s	s += a[i]
b000	1A	doit	nop		
4002	1B		store	s	
3001	1C		load	i	i += 1
5006	1D		add	one	
4001	1E		store	i	
C011	1F		jmp	loop	goto loop
3002	20	xit	load	s	print s
1800	21		printH		
0000	22		halt		

Listing 1.9: IBCM array index program

A quick note before the full analysis. Notice that addresses 08 and 09 are blank – this was done to leave space for additional variables. If one has to add a variable into the program later, shifting all of the successive instructions down is a frustrating task, as all the addresses (of the jump targets, etc.) need to be shifted as well. So we left a few blank lines in case we needed additional variables later on.

The challenging part of this program is the array subscripting. In our pseudo code – and in most pro-

gramming languages – we use a syntax such as `a[i]`. But the IBCM does not have any array subscripting instruction – indeed, it cannot, as that requires two values (the array base and the index). Thus, we have to *build* the instruction to execute.

Our goal is to load the current sum (of the array elements processed so far) into the accumulator, execute the instruction to add the current `a[i]` to the accumulator, and store that back into the sum variable. This is done in instructions 19, 1a, and 1b: instruction 19 loads the current sum into the accumulator, instruction 1a is our special instruction that adds `a[i]` to the accumulator (how this works is next), and instruction 1b stores the updated result back into the sum variable (address 002).

Thus, we want to create an instruction that will add a given `a[i]` to the accumulator. So we know it will be an add instruction (opcode of 5). The address that we want to add to the accumulator is the array base plus the index.

For example, if the array starts at address 100, and we are trying to add the value at index 5 of the array, our instruction would have opcode of 5 (an add instruction), and address 105 (100 for the array start plus 5 for the current index). Thus, our instruction would be 5105.

To create this during the program execution, we start with 5000, which sets the opcode of the instruction to the correct value for an add instruction. This is done in instruction 15, which is a load of the `adit` variable (address 007), which has value 5000. To this value of 5000 we add the array base (instruction 16) and then the index (instruction 17). Instruction 18 then stores that instruction at the appropriate address (address 1a), overwriting the `nop` that is there with our updated instruction. Once the current sum is loaded (instruction 19), our custom instruction is executed (instruction 1a), followed by the sum being stored back into memory (instruction 1b).

1.7.3 Recursive multiplication

The next program computes the product of two numbers, x and y, through a recursive multiplication subroutine that uses only addition. Both x and y are read from the keyboard, and the resulting product is printed to the screen. The IBCM version is just over 100 lines of opcodes.

The C++ code for the multiplication program can be seen in Listing 1.10. We did not create a tail recursive `multiply()` routine, as the IBCM compiler and language is (intentionally) far too simple to optimize for tail recursion.

This IBCM program creates a stack similar to x86: the stack starts at the end of addressable memory, and grows downward. An activation record is created for each recursive call, which consists of the two parameters and the return address – other fields typically in an activation record (e.g., backup of registers) are not necessary in IBCM. The `brl` instruction was used to allow for subroutine calls – the return address is saved in the accumulator, and is stored on the stack immediately upon subroutine activation. We were able to run

```
int mult(int x, int y) {
  if ( y == 0 )
    return 0;
  else
    return x+mult(x,y−1);
}


int main(void) {
  int x, y;
  cin >> x;
  cin >> y;
  cout << mult(y,x) << endl;
}
```

Listing 1.10: C++ multiplication program

the program with the second parameter (which is decremented in each recursive call) set as high as 1,243 (0x4db), beyond which point IBCM runs out of available memory for the stack.

This recursive multiplication program is beyond what we would expect of a student to be able to program after a one-week introduction to IBCM. However, it is very illustrative of two important points about IBCM. One is that complex functionalities (such as multiplication) can be achieved by using only the simple capabilities available in IBCM (such as addition). The other is that subroutines are fully realizable in IBCM.

This program is shown in Listings 1.11 and 1.12. Note that the only part of the program that the simulators read in is the first four characters on the line. The rest of the line in the text file is solely for comments. The column headers and extra blank lines for comments are shown in the diagram for clarity, but would not be included in the IBCM source code file.

1.8   Turing Completeness

We were conflicted as to how much to discuss about Turing completeness in this chapter. IBCM is similar to a Random Access Stored Program (RASP) machine [15], which is itself Turing complete, so it is perhaps not surprising that IBCM is also Turing complete. However, we felt that breaking down a complex task (a Turing machine simulator) and programming it into IBCM was a worthwhile task to discuss, as it emphasizes a primary design goal of IBCM – that you can take any algorithm and write it in IBCM.

Obviously, no physical computer with a finite amount of memory can be truly Turing complete. Thus, we will instead show that the IBCM computational model is Turing complete.

We define the *IBCM computational model* as the same IBCM computer defined above, but allowing any sized integer to be held in a single memory location, as well as having an infinite amount of memory. Thus, other fields that make up part of a given instruction, such as the 'address' or 'count' fields (see Figure 1.2), can also hold any size integer.

Given this model of computation, we will show how to simulate a Turing machine in IBCM. Hopcroft and Ullman [4] define a Turing machine as a 7-tuple $M = \langle Q, \Gamma, b, \Sigma, \delta, q_0, F \rangle$ where:

- Q is a finite set of states
- Γ is the finite set of allowable tape symbols
- $b \in \Gamma$ is the blank symbol
- $\Sigma \subseteq \Gamma \setminus b$ is the set of input symbols
- $\delta : Q \times \Gamma \to Q \times \Gamma \times \{L, R\}$ is the transition function
- $q_0 \in Q$ is the initial state
- $F \subseteq Q$ is the set of final states

We will make two modifications to the above definition, to allow for ease of implementation in IBCM. We will allow a no-shift transition of S (in addition to L and R), which will not move the tape. The Turing machine programs allowed by the no-shift are equivalent to the ones described above [16]. Furthermore, we will define f, a single final state, which all states in F move to on a no-shift transition.

To represent the transition functions in IBCM, we will represent state $q \in Q$ and symbol $\Sigma \in \Gamma$ each as a (16-bit) word in IBCM. Thus, states and symbols will be a single integer each. While this limits the number of states (and symbols) to $2^{16} = 65,536$ in our IBCM implementation, it is not limited in the formal IBCM computational model. Thus, one can encode any amount of states, symbols, and transition functions into IBCM's memory.

Mem	Loc'n	Label	Opcode	Addr	Comments
					push the return address onto stack
4002	27	mult	store	tmp	store the return address into tmp
3001	28		load	sptr	load the stack pointer
5004	29		add	store	create a store instruction
402c	2a		store	pos3	store the store inst into pos3
3002	2b		load	tmp	load up the value to push onto stack
b000	2c	pos3	nop		will hold the push−to−stack inst
3001	2d		load	sptr	load the stack pointer
6007	2e		sub	one	decrement it
4001	2f		store	sptr	store it back
					set possible return value of zero
3006	30		load	zero	load zero into the accumulator
400a	31		store	retval	if we are returning, store zero in the return value
					get the 2nd parameter, compare to 0
3001	32		load	sptr	load the stack pointer
5003	33		add	load	create a load instruction
5009	34		add	three	move to the pos of the 2nd parameter
4036	35		store	pos4	store the load instruction into pos4
b000	36	pos4	nop		will hold inst to load 2nd parameter
4002	37		store	tmp	store it in tmp for use on line 049
d03a	38		jmpe	ret	if 0, we're done with the recursion
c046	39		jmp	recurse	otherwise call ourselves recursively
					in order to return, we need to create a jump inst, clean up the stack, and return 0
3001	3a	ret	load	sptr	load the stack pointer
5003	3b		add	load	create a load instruction
5007	3c		add	one	move to the pos of return address
403e	3d		store	pos5	store the load instruction into pos5
b000	3e	pos5	nop		will hold inst to load return addr
5005	3f		add	jmp	create a jump instruction
4045	40		store	pos6	store the jmp instruction into pos6
3001	41		load	sptr	load the stack pointer
5007	42		add	one	add one to 'remove' the return addr
4001	43		store	sptr	store it back
300a	44		load	retval	load the return value
b000	45	pos6	nop		will hold the jmp (return) inst;
					to call ourselves recursively, we push the 2 parameters onto the stack stack, then call with brl.
					load 2nd parameter, decrement it, then push it onto the stack
3001	46	recurse	load	sptr	load the stack pointer
5004	47		add	store	create a store instruction

Listing 1.11: IBCM multiplication program, part 1

Mem	Loc'n	Label	Opcode	Addr	Comments
404b	48		store	pos7	store the store instruction in pos2
3002	49		load	tmp	the 2nd parameter is already in tmp
6007	4a		sub	one	subtract one
b000	4b	pos7	nop		will hold the push−to−stack inst
3001	4c		load	sptr	load the stack pointer
6007	4d		sub	one	decrement it
4001	4e		store	sptr	store it back
					load the 1st parameter into tmp
3001	4f		load	sptr	load the stack pointer
5003	50		add	load	create a load instruction
5009	51		add	three	move to position of 1st parameter
4053	52		store	pos8	store the load instruction into pos4
b000	53	pos8	nop		will hold inst to load 1st parameter
4002	54		store	tmp	store it in tmp for use on line 058
					push 1st parameter onto stack second
3001	55		load	sptr	load the stack pointer
5004	56		add	store	create a store instruction
4059	57		store	pos9	store the store inst into pos1
3002	58		load	tmp	load value to push onto the stack
b000	59	pos9	nop		will hold the push−to−stack inst
3001	5a		load	sptr	load the stack pointer
6007	5b		sub	one	decrement it
4001	5c		store	sptr	store it back
					call the subroutine
f027	5d		brl	mult	using the branch−and−link inst
					take the return value, and add 1st
					parameter to it, and return that
4002	5e		store	tmp	store the return value in tmp
3001	5f		load	sptr	load the stack pointer
5003	60		add	load	create a load instruction
5007	61		add	one	move to the position of the 1st
					parameter (one because the ret
					addr is not on the stack)
4063	62		store	pos10	store the load inst into pos4
b000	63	pos10	nop		will hold the instruction to load
					the 1st parameter
5002	64		add	tmp	add returned value to the 1st param
400a	65		store	retval	store it for later
					clean up stack (pop the two params)
3001	66		load	sptr	load the stack pointer
5008	67		add	two	increment the stack pointer by two
					for the two parameters
4001	68		store	sptr	store it back
c03a	69		jmp	ret	jump to the return part of this code
					(return value in retval)

Listing 1.12: IBCM multiplication program, part 2

Figure 1.9: Four state Busy Beaver automaton

To simulate a Turing machine, we will define an arbitrary memory location to represent the current state of the Turing machine, and another arbitrary memory location to contain the current address of the head of the tape. The transition function quintuples, δ, will start at a specific (but arbitrary) memory address, take up five words each, and will contain the five parts listed above ($Q, \Gamma, Q, \Gamma, \{L, R, S\}$). The tape itself will start at different arbitrary memory location. Furthermore, we define a initial state q_0 and a (single) final state f. The blank symbol b will be an arbitrary value, such as -1 (0xffff in 16-bit 2's complement integer).

Any Turing machine that requires a significant amount of tape will need to be a one-way tape Turing machine, as the program code and state transitions will lie at a lower address than the initial head position. Thus, only a finite amount of tape space is available in the lower memory address direction. The particular automaton example that we provide, below, uses a two-way tape, but that is because we know the finite amount of tape space necessary. Note that one-way tape Turing machines are equivalent in computational power to two-way tape Turing machines [4].

What is needed, then, is an IBCM program that will iterate through the following steps:

- Read the current state s, initially set to the start state. If the current state is the (single) final state, then exit.
- Read the current head position.
- Read the current input symbol t at the head position
- Search the list of transition functions until the appropriate one is found, based on the current state s and the input symbol t.
- Perform the action specified in the transition function by updating the state s, writing the specified symbol to the tape position, and then moving the tape left or right (or, on an S, not at all).

We have developed such a program, described here. The full listing of the program is available online [2]. The program consists of 67 IBCM commands, and 15 variables – note that numeric constants are considered variables in an IBCM program. This program only used half of the instructions: `halt`, `load`, `store`, `add`, `sub`, `nop`, `jmp`, and `jmpe`.

To test the Turing machine, we choose a four state Busy Beaver automaton, which is described in more detail in the Wikipedia page on Turing machine examples [17]. The Mealy machine finite automaton is shown in Figure 1.9. For each transition, the input symbol (0 or 1) is shown, along with the tape direction to move (L, R, or S). Note that in this automaton, upon each transition a 1 is written to the output tape; this is not shown in the figure to improve clarity. Also recall that the S transition means to not move the tape, and is used only on the transition to the final state.

Our implementation can utilize a two-way tape, although the tape in one direction is finite. We define memory address 1 as the current state variable, and address 2 to store the location of the current head

position. The transitions start at address 0x060, as our program takes up 82 (or 0x052) instructions. The states are numbered as per the diagram in Figure 1.9, with 0 being the initial state and 3 being the final state. The program has constants that specify both the initial and final states of the automaton.

The encoding of the automaton shown in Figure 1.9 is very straightforward. The transition from 2 to 1, executed on an input symbol of 0, will print 1 to the tape and then move the tape to the left. The quintuple to be encoded is $Q \times \Gamma \to Q \times \Gamma \times \{L, R, S\}$. The respective values for this transition are $2, 0, 1, 1, 0$; we map the integer values $\{0, 1, 2\}$ to, respectively, the transition directions $\{L, R, S\}$.

1.9   Emulating IBCM in C++

How might one write software to emulate an IBCM machine in C++? A switch statement with 16 cases, perhaps. But how to decode the instructions?

Let's assume we had to write a C++ program that could extract the parts of an IBCM instruction. How to do it? Assume the instruction is in an `unsigned int x`. One way to decode it is shown in Listing 1.13

```
unsigned int opcode = (x >> 12) & 0x000f
unsigned int ioshiftop = (x >> 10) & 0x0003
unsigned int address = x & 0x0fff
unsigned int shiftcount = x & 0x000f
```

Listing 1.13: Decoding an IBCM instruction in C++

What about encoding? Assuming we have (unsigned ints) opcode, ioshiftop and shiftcount, the decoding is shown in Listing 1.14

```
unsigned int instruction = (opcode << 12) | (ioshiftop << 10) | shiftcount
```

Listing 1.14: Encoding an IBCM instruction in C++

This ends up being a rather frustrating program to write. If the instruction set being dealt with is more complicated than IBCM, as is the case with x86 or MIPS instructions, then the above is very difficult to do without any errors.

Listing 1.15 shows a data structure to make it easier. While the IBCM is a big-Endian machine, the following code may be running on a big or little-Endian machine. The `BIG_ENDIAN` and `LITTLE_ENDIAN` defines specify the Endianness of the host machine.

We would use the data structure as shown in Listing 1.16.

1.10   Pedagogy

IBCM was originally developed at the University of Virginia to complement our CS 3 course entitled *Program and Data Representation*, which is still taught today. This course shows how one represents both data and program code from high levels – such as abstract data types – all the way down to the lowest (software) level, which is the IBCM machine language.

IBCM allows students to easily make the mental connection between assembly opcodes and the machine language that they get translated into. At the University of Virginia, we follow the presentation of

```
union ibcm_instruction {
#ifdef BIG_ENDIAN // the IBCM is big endian
    struct { unsigned char high, low; } bytes;
#else
#ifdef LITTLE_ENDIAN
    struct { unsigned char low, high; } bytes;
#else
#error Must define BIG_ENDIAN or LITTLE_ENDIAN
#endif // LITTLE_ENDIAN
#endif // BIG_ENDIAN
    struct { unsigned int op:4, unused:12; } halt;
    struct { unsigned int op:4, ioopt:2, unused:10 } io;
    struct { unsigned int op:4, shiftop: 2,
        unused:5, shiftcout:5; } shifts;
    struct { unsigned int op:4, address:12; } others;
};
```

Listing 1.15: C++ data structure to ease IBCM instruction encoding

```
// read in instruction into unsgined chars a and b
ibcm_instruction inst;
inst.high = a;
inst.low = b;
if ( inst.halt.op == 0 ) { // halt
} else if ( inst.io.op == 1 ) { // io
        cout << inst.io.ioopt << endl;
} else if ( inst.shifts.op == 2 ) { // shifts
        cout << inst.shifts.shiftop << endl;
        cout << inst.shifts.shiftcount << endl;
} else { // all others
        cout << inst.others.address << endl;
}
```

Listing 1.16: Using the C++ data structure to encode IBCM instructions

IBCM with a two-week introduction to x86 assembly language. This allows students to understand both machine language, as well as a modern processor's assembly language (we use Intel x86), without having to delve into the details of x86 machine language.

A specific design decision with IBCM was to include only basic operations – for example, multiplication and division are not included, but can be replicated by using repeated addition or subtraction. We wanted students to see that any program, no matter how complicated, can be broken down into very simple instructions. Indeed, this is the point of the concept of Turing machines, but they are often not taught when students are first seeing machine language.

Students are exposed to a number of concepts during the IBCM module that they often have not seen previously. They become aware that in both assembly language and machine language, data is untyped, and the operations on the data determine the type – this is quite different than the typed high-

level programming languages to which they are accustomed. By this point in our course, students have been exposed to how a 32-bit value can be interpreted either as a two's-complement signed integer or an IEEE 754 floating point number.

Another concept taught through IBCM is self-modifying programs. A non-trivial IBCM program requires arithmetic on instructions – in fact, only the first example program shown in Listing 1.7 did not use this feature. Array indexing, for example, requires starting with a `load` or `add` instruction, and adding to that value both the base address of the array and the current index. This value is then stored in a memory address which is shortly thereafter executed. The second example program provided to the students uses this feature. While many systems explicitly try to prevent self-modifying code, as that is an exploit used by a significant amount of malware, it is still a concept that the students should be familiar with.

The development of self-modifying IBCM programs leads to another pedagogical goal of the course: the interplay between data and program code. Indeed, there is little difference between data and program code, other than the values (obviously), and how it is interpreted (and, in modern systems, what segment of an executable in which the data is found). This concept seems trivial to instructors, but is one that students who have only programmed using high-level languages are often unfamiliar with.

What IBCM does not teach, of course, is how binary machine language instructions are executed on the processor. Understanding of this material is typically beyond all but senior-level undergraduate courses; at many institutions, this is also outside the standard computing curriculum.

1.10.1   Materials Available

We have developed and made available a wide range of materials for the purpose of teaching the IBCM module. The materials are all released under various Creative Commons licenses. They are available online [2], and consist of:

- A PowerPoint slide set to introduce the concepts during a lecture-based course. This 42 slide set takes about three 50-minute lectures to present.
- A *Principles of Operation* document, which describes the IBCM computer and language, and how to write a program. It covers similar content to the lecture slides.
- Sample programs, one of which was shown in Listing 1.7, above. We also provide sample programs on array indexing, for example. All the programs mentioned in this article are available online.
- Sample student assignments, which require students to write additional IBCM programs beyond the sample programs provided. One of the assignments is to write a quine, or an IBCM program that will print itself out – the smallest quine produced is nine IBCM commands.
- An online PHP/Javascript simulator, which is the primary way that the students program in IBCM. This is shown above in Figure 1.8. The PHP is used to allow loading of an IBCM program from a text file; the Javascript implements the IBCM simulator in the browser itself. The simulator works across all major browsers on all major operating systems.
- A C++-based command-line program, which can both compile and execute IBCM programs. Not surprisingly, this is much faster than the online simulator. This is particularly useful for automated compilation and execution of IBCM programs for grading, or for very long programs, such as the recursive multiply routine described above.

Furthermore, a GUI-based tool for executing IBCM programs is available separately [14]. This tool allows for drag-and-drop loading of IBCM files into the GUI, and compiles natively for each operating system.

We very specifically have not developed an IBCM assembler, which would take in the opcodes in an assembly language format and output hexadecimal machine code. The purpose of IBCM is to teach the students machine language; given an IBCM assembler, this module ends up being just a different assembly language for the students to learn.

1.10.2 Related Work

We are certainly not the first to propose a simplified machine language as a pedagogical tool. Andrew Tanenbaum's original 1984 text, *Structured Computer Organization* – now in its 5th edition – presents the Mic-1 micro architecture and the Mac-1 machine language [11]. Indeed, the Mac-1 machine language has many similarities with IBCM – this is not surprising, as there are many common elements that must be present in all small instruction set machine languages. Further research in that decade presented implementations of those languages [6, 7]. The Mic-1 simulator, which is focused on assembly language, is available online [10]. Simulators for the Mac-1 machine language do exist, but seem to be independently developed, and without a modern set of implementation software.

A number of high quality simulators exist at the assembly level. Pep/8 [12, 13] is a 16-bit CISC architecture designed to teach assembly language concepts. While it can be used for machine language – and can trace program execution at the machine language level – the primary pedagogical design is at the assembly language level. Another example is SPIM, which is a full featured MIPS 32-bit simulator [5].

Additional research on machine language simulators has often focused on the register transfer level [8], or is restricted to a single client operating system [1].

More recent research by Stone has focused on a similar machine language implementation [9]. Although developed independently, our research can be seen as an extension of Stone's, as we add a number of additional aspects: significant pedagogical tools, a fully downloadable package so this system can be used in any course, a proof of Turing completeness, and a discussion of pedagogical concerns.

We are not aware of any machine language simulators that are available with the set of modern tools that we present with IBCM.

1.10.3 Results

At the University of Virginia, this module is taught about half-way through our CS3 course, which is where we teach data structures. Our CS1 and CS2 course are both in Java. At this point in the CS3 course, the students have learned to implement a number of data structures in C++. We introduce machine language using IBCM for one week, follow that with two weeks of assembly programming (x86), and then return to C++ for the remainder of the semester. The lectures used to teach IBCM typically take three 50-minute class periods. This is followed by an IBCM lab the following week. All of the lecture slides and labs are available online [2].

We have taught machine language using IBCM in our CS3 class for over a decade here at the University of Virginia. Over the years, student reactions to IBCM have varied greatly. With the usage of the modern tool set presented in this article, those reactions have generally been positive, as shown below While they often do not like being constrained to such a limited set of instructions, they see the purpose of learning machine language, and generally enjoy the IBCM assignments. We have found that the quality of the software tools is directly related to student perception of IBCM – in the past, when our IBCM simulator was less refined, student reaction was significantly more negative.

Objective comparison with other programming languages, including assembly language, are difficult due to the vast differences between both the capabilities and the required learning curve. We have instead focused on subjective assessment. In the most recent semester in which IBCM was used (fall of 2010), six questions were asked of the students. All questions were asked on a Likert scale, where 1 means strongly agree, 2 is agree, 3 is neutral, 4 is disagree, and 5 is strongly disagree. For all the questions, $n = 89$.

The first two questions focused on how much the students felt they learned from the IBCM module. The middle two questions focused on the ease of use and enjoyability. And the last two questions focused on how worthwhile this module was, both for this course and future courses. The results are shown in Table 1.4.

Question	Avg	Stdev
IBCM increased my understanding of the basics of machine language	1.67	0.58
IBCM increased my understanding of how computers work at a low level	1.89	0.65
IBCM was easy to use, once I got the hang of programming in it	1.92	0.88
I enjoyed learning IBCM	2.01	0.98
Considering what was taught, IBCM was a worthwhile module to have in this course	1.75	0.68
IBCM should be used in future iterations of this course	1.76	0.72

Table 1.4: IBCM survey results

The results clearly show that the module was generally well received. The lowest score, enjoyment of learning IBCM, still was in the 'agree' category.

1.11   Conclusions

We have used IBCM for over a decade at the University of Virginia in our CS3 class. During that time, we have refined it into the pedagogical tool presented here. With the set of current pedagogical tools provided, we have found it to be an effective means of teaching the basic concepts of machine language without having to go into the complexity of modern machine languages that is beyond the scope of a lower-level undergraduate course. All of the necessary materials are available online for adoption at other institutions.

Bibliography

[1] B. Lewis Barnett, III. A visual simulator for a simple machine and assembly language. In *SIGCSE '95: Proceedings of the twenty-sixth SIGCSE technical symposium on Computer science education*, pages 233–237, New York, NY, USA, 1995. ACM.

[2] Aaron Bloomfield. Itty bitty computing machine online. `http://www.cs.virginia.edu/~asb/ibcm/`.

[3] J.F. Cleverley and D.C. Phillips. *Visions of childhood: Influential models from Locke to Spock*. Teachers College Press New York, 1986.

[4] John E. Hopcroft, Rajeev Motwani, and Jeffrey D. Ullman. *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley Publishing Company, 3rd edition, 2006.

[5] James Larus. SPIM: A MIPS32 Simulator. `http://pages.cs.wisc.edu/~larus/spim.html`.

[6] Jerry E. Sayers and David E. Martin. A hypothetical computer to simulate microprogramming and conventional machine language. *SIGCSE Bull.*, 20(4):43–49, 1988.

[7] Delmar E. Searls. An integrated hardware simulator. *SIGCSE Bull.*, 25(2):24–28, 1993.

[8] Dale Skrien and John Hosack. A multilevel simulator at the register transfer level for use in an introductory machine organization class. *SIGCSE Bull.*, 23(1):347–351, 1991.

[9] Jeffrey A. Stone. Using a machine language simulator to teach cs1 concepts. *SIGCSE Bull.*, 38(4):43–45, 2006.

[10] Andrew Tanenbaum. Mic-1 download website. `http://www.ontko.com/mic1/`.

[11] Andrew S. Tanenbaum. *Structured Computer Organization*. Prentice Hall PTR, Upper Saddle River, NJ, USA, 1984.

[12] J. Stanley Warford and Chris Dimpfl. The pep/8 memory tracer: visualizing activation records on the run-time stack. In *Proceedings of the 41st ACM technical symposium on Computer science education*, SIGCSE '10, pages 371–375, New York, NY, USA, 2010. ACM.

[13] J.S. Warford. *Computer Systems*. Jones & Bartlett Learning, 2009.

[14] Jacob Welsh. JWelsh's useful programs. `http://www.eemta.org/~jwelsh/progs/`.

[15] Wikipedia. Random access stored program machine. `http://en.wikipedia.org/wiki/Random_access_stored_program_machine`.

[16] Wikipedia. Turing machine. `http://en.wikipedia.org/wiki/Turing_machine`.

[17] Wikipedia. Turing machine examples. `http://en.wikipedia.org/wiki/Turing_machine_examples`.